

50 Things You Could Be Automating With Your AI Agent

The playbook for turning invisible manual work into agent-run systems: sales, marketing, support, ops, engineering — and you.

Every one of them comes with a receipt.

An automation you can't check isn't an automation. It's a rumour.

I have spent the last year building teams of AI agents and measuring what they actually cost. The most expensive lesson had nothing to do with money.

I once shipped a health check for a fleet of agents. It watched them, and it told me they were alive. It said so four separate times. They were dead. The check was matching against its own process – it was finding *itself* and reporting a heartbeat. My monitoring didn't miss the incident. My monitoring **was** the incident.

That is the whole risk of automating your company. Not that the robot does the wrong thing loudly – you'd catch that by Tuesday. It's that it does *nothing*, quietly, and everything looks fine. A queue that is wrong looks exactly like a queue that is fine.

So this isn't a list of prompts. It's a list of **jobs** – and every single one of them ships with a *receipt*: the thing you look at to know it really ran. If you take nothing else from this guide, take the receipts.

Rule 1 – Start where it's boring, repetitive, and checkable.

Not where it's hardest. Not where it's most impressive in a demo. The first automation should be something you already do by hand every week and could verify in ten seconds. Boring work is where agents are superhuman and where a mistake costs you nothing.

Rule 2 – If you can't name the receipt, don't build it.

Before you automate anything, finish this sentence: *"I'll know it worked because I'll see ____."* A message in a channel. A row in a sheet. A link to the thing it made. If you can't finish that sentence, you're not building an automation – you're building a place for work to disappear.

Rule 3 – Silence must be loud.

The default state of a broken automation is silence, and silence feels like peace. Make the absence of work an event: *"O invoices chased"* on a day with overdue invoices should page you as hard as an error would. Every serious failure I've had was something that quietly did nothing and looked fine.

Pick **three** from the next pages. Not thirty. Three you already do by hand, that you'd notice within a day if they stopped. Run them for two weeks. Then come back.

01 SALES & PIPELINE

The work that happens *around* selling is the work that doesn't get done.
None of this replaces the call. All of it replaces the hour after it.

01 Turn every sales call into a CRM update

It listens to the recording, writes the summary, updates the deal stage, and lists what you promised the customer.

RECEIPT The deal record changed today. If you had a call and nothing moved, that's a bug – not a quiet day.

02 Draft the follow-up while the call is still warm

Within ten minutes of hanging up, a draft is sitting in your inbox with what you agreed and the next date. You edit and send.

RECEIPT A draft exists before you've made coffee. No draft = no call happened, or the agent is down.

03 Tell you which deals have gone quiet

It watches for the deals where nobody has spoken in fourteen days and puts them in front of you, oldest first.

RECEIPT A weekly list with names on it. An empty list on a week you know is messy means it isn't looking.

04 Prep you before the intro call

Fifteen minutes before the meeting: who they are, what they raised, what they shipped last month, and the three questions worth asking.

RECEIPT A one-pager in your inbox with a timestamp before the meeting starts. Late = useless.

05 Keep the pipeline honest

It sweeps for deals rotting in a stage nobody has touched, and asks the owner one question: is this real?

RECEIPT The stale count, posted every Monday. If it's ever zero, be suspicious – pipelines are never that clean.

06 Answer the security questionnaire

It fills the 200-row spreadsheet from your previous answers and flags only the rows it genuinely doesn't know.

RECEIPT The count of rows it left blank. A confident 0-blanks answer is the one to distrust – go read it.

07 Write the pipeline note for your board

It turns the CRM into three paragraphs a human investor would actually read, with the numbers attached.

RECEIPT It's in your drafts before the board email is due, with every number traceable back to a deal.

The trap here is asking an agent to have opinions. Don't. Ask it to *carry* yours – into more places, more often, from things you already made.

08 Turn one long thing into ten short things

A talk, a doc, a customer call – cut into posts, each one able to stand alone.

RECEIPT Ten drafts land in a folder. You approve or kill each one. Nothing publishes without a click.

09 Write from what you actually shipped

It reads the week's changelog and drafts the customer-facing note: what changed, why you'd care.

RECEIPT A draft per release. If you shipped and there's no draft, the pipe is broken.

10 Keep the swipe file

Everything that performed well – yours and other people's – captured with a line on why it worked.

RECEIPT The file grows every week. A file that stopped growing is a dead agent.

11 Find the conversations worth joining

It surfaces the threads where someone is asking, right now, the exact question your product answers.

RECEIPT A short list, daily, with links. You reply as yourself. The agent never replies.

12 Refresh the page that's gone stale

It flags copy that promises things you no longer do, or misses things you now do.

RECEIPT A diff you can read in a minute. Approve, edit, or ignore.

13 Turn real customer questions into content briefs

It reads support tickets, clusters the repeated questions, and hands you the ten pieces worth writing.

RECEIPT Each brief cites the tickets it came from. A brief with no tickets under it is invented – bin it.

14 Send the 'what shipped' digest

Monthly, to customers, in your voice – every item linked to the thing itself.

RECEIPT It goes to your drafts, not to the customers. You press send. Always you.

Keep a human on the send button for anything a customer reads. Automate everything that happens *before* that button.

15 Triage by what people mean, not what they typed

Bug, billing, feature request, angry – sorted before anyone opens the inbox.

RECEIPT Every ticket carries a label and a confidence. Low-confidence ones go to a human queue, visibly.

16 Draft the answer, with the doc link attached

The reply is written and the relevant doc is found. A human reads it and hits send.

RECEIPT Time-to-first-draft. If a ticket sits with no draft, the agent didn't see it.

17 Turn the question you keep answering into a doc

The fifth time the same question arrives, it writes the doc and tells you.

RECEIPT A new doc appears, citing the five tickets. That count is the whole justification.

18 Spot the customer who's about to leave

Logins fall off a cliff, usage halves, the champion goes quiet – you hear about it this week, not at renewal.

RECEIPT An alert with the specific numbers that moved. No numbers, no alert – hunches don't count.

19 Nudge the customer who never finished setup

It notices who signed up and never did the one thing that makes the product work, and drafts a human nudge.

RECEIPT The list of who was nudged, and who converted after. If nobody converts, kill it – it's spam.

20 Turn a vague bug report into something reproducible

"It's broken" becomes steps, environment, and the log line that matters.

RECEIPT The issue it files either reproduces or it doesn't. That's a hard, honest scoreboard.

21 Write the refund decision memo

It assembles the account history, the policy, and a recommendation. A human decides.

RECEIPT A memo per request. The decision line is blank until a person fills it in – by design.

One hard rule, and I mean it: **an agent may draft a money decision, never make one.** Everything below stops one step short of the transaction, on purpose.

22 Reconcile the numbers you keep eyeballing

The two spreadsheets that should agree, checked every night, with only the differences shown.

RECEIPT A nightly "matched" or "3 rows differ." A night with no message at all means it didn't run – treat that as a failure.

23 Chase the invoices nobody wants to chase

Overdue by seven days, it drafts the polite nudge. Overdue by thirty, it drafts the less polite one.

RECEIPT The list of who it chased, posted where you'll see it. Zero chased on a day with overdue invoices is a bug.

24 Audit the subscriptions you forgot you had

Every recurring charge, matched against whether anyone logged in this quarter.

RECEIPT A list with a total at the bottom. Most founders find four figures a year on the first run.

25 Watch the cloud bill for the shape change

Not the total – the shape. It tells you when something that cost \$8 a day starts costing \$80.

RECEIPT An alert naming the line item and the day it changed. Totals hide this; line items don't.

26 Turn receipts into a ledger

Photograph, forward, or drop it. It's categorised and filed before your accountant asks.

RECEIPT The count filed this month vs. the count on your card statement. Those numbers must match.

27 Read the contract first

It flags auto-renewal, liability, exclusivity, and anything unusual – with the clause quoted.

RECEIPT Flags always quote the clause. A flag with no quote is a hallucination, and you'll spot it instantly.

28 Collect the compliance evidence continuously

The screenshots and logs your auditor will want, gathered as you go instead of in one miserable week.

RECEIPT The evidence folder has this week's date in it. Stale folder = dead agent.

29 Write the Monday numbers email

The five numbers that matter, what changed, and one sentence on why – before you're awake.

RECEIPT It arrives every Monday. A Monday with no email is an incident, not a break.

Even if you don't write code, read this one. It's where your team is quietly losing days – and where the receipts are easiest to see.

30 **Review the code before a human does**

First pass on every pull request. It only speaks when it can point at a line and say why.

RECEIPT Comments cite file and line. A review with no line numbers is noise – configure it to stay quiet instead.

31 **Write the test that reproduces the bug**

Before anyone fixes anything, there's a failing test that proves the bug is real.

RECEIPT The test fails before the fix and passes after. If it passes before the fix, it's testing nothing.

32 **Keep dependencies moving**

Small upgrades, continuously, with the test suite as the judge – instead of one terrifying upgrade a year.

RECEIPT Green suite, merged. Red suite, it stops and tells you which package broke.

33 **Draft the postmortem while it's fresh**

It assembles the timeline from the alerts, the deploys and the chat, so a human writes the analysis, not the archaeology.

RECEIPT A timeline with timestamps you can check against the logs yourself.

34 **Catch the slow query before your customers do**

The database call that runs 400 times when it should run once, flagged in review.

RECEIPT It names the query and the count. Numbers, not adjectives.

35 **Keep the docs honest**

When the code changes and the README doesn't, it notices and drafts the fix.

RECEIPT The diff between what the docs claim and what the code does. Empty diff = they agree.

36 **Check the migration won't take the site down**

The schema change that locks a table for nine minutes on production data, caught before it ships.

RECEIPT A verdict with the estimated lock time, on a copy of real data volumes – not on your laptop's 40 rows.

37 **Hunt the flaky test**

It runs the suspicious test a hundred times and tells you the failure rate, instead of everyone shrugging and hitting retry.

RECEIPT "Fails 6 times in 100." Now it's a number, and numbers get fixed.

38 **Build the prototype so the meeting has something to look at**

A clickable version of the idea, same day, before anyone commits a quarter to it.

RECEIPT A link that opens. Every button goes somewhere, or it isn't done.

39 Audit the codebase you're about to inherit

Before you hire, before you acquire, before you say yes – what's actually in there.

RECEIPT A ranked report where every finding points at a file. Findings without files are vibes.

40 Turn the error tracker into a to-do list

Ranked by how many real customers hit it, not by how loud it is.

RECEIPT The rank order and the user counts behind it. You'll disagree with the top item once, and be wrong.

The line here is sharp: an agent may *prepare* a judgement about a person. It may never *make* one. Prepare, then you decide.

41 Screen applications against the actual job

Not keywords – the job. It sorts, explains why, and never rejects anyone by itself.

RECEIPT Every ranking has a written reason you can argue with. If you can't argue with it, it isn't a reason.

42 Turn interview notes into a scorecard

Your scribbles, structured against the same criteria for every candidate, so you compare like with like.

RECEIPT A filled scorecard per interview, within the hour. Empty fields stay visibly empty.

43 Get the reply out the same day

Offer or rejection, drafted immediately, so nobody waits two weeks to hear no. You send it.

RECEIPT Time from decision to draft. The number that matters – and the one that makes people respect you.

44 Build the 30-day plan before they start

Per role: what they read, who they meet, what they ship in week one.

RECEIPT The plan exists before day one. A plan written on day three is a plan that doesn't exist.

45 Prep your 1:1s

What they said last time, what they were blocked on, whether it moved.

RECEIPT It shows you last week's blocker and whether it's still there. That one line changes the meeting.

The last five are the ones founders skip and then wish they hadn't. Your attention is the scarcest thing in the company. Defend it like a budget.

46 Write the Friday brief you never write

What moved, what didn't, what you said you'd do and didn't do.

RECEIPT It arrives Friday. The "didn't do" section is the one worth reading, so it never gets dropped.

47 Kill the meeting that should have been a document

It drafts the doc from the agenda. If the doc answers it, cancel the meeting.

RECEIPT Meetings cancelled. Count them. That's hours back, and it's a real number.

48 Triage your inbox by 'does this need a founder?'

Most of it doesn't. It surfaces the handful that genuinely need you.

RECEIPT The count it held back, visible daily. If you never disagree with what it filtered, you aren't checking.

49 Research the decision before you make it

The memo – options, trade-offs, what it costs – sitting there before the call, not improvised during it.

RECEIPT Every claim in the memo carries a source. A claim without a source is an opinion wearing a suit.

50 Watch your other agents

One agent whose only job is to notice when another one has quietly stopped doing anything.

RECEIPT It reports on the ones that produced nothing. Then break it on purpose and confirm it screams. If it doesn't scream, you are back where you started – and you won't know it.

Where this goes wrong, and how you'll know.

"The first thing that breaks won't be the agent. It'll be your ability to tell whether the agent did anything at all."

Here's the failure I want you to expect, because it's the one I keep having. You wire something up. It works. You stop watching it. Three weeks later you find out it stopped working on day four – and nothing anywhere told you, because *nothing* is exactly what a broken automation produces.

Every serious failure in my own systems has had that same shape. A monitor that matched itself and reported health. A publish button that reported success and posted nothing. A guard against double-posting that never ran, so it looked like it worked. All of them were **instruments that couldn't tell success from nothing**.

Which is why the fix is boringly mechanical, and why I'll keep repeating it: **test the refusals first**. Before you trust an automation, break it on purpose. Unplug it and confirm something screams. If nothing screams, you don't have an automation – you have a rumour that one exists.

Your first two weeks

Week 1. Pick three from this list. Boring, repetitive, checkable. Write down the receipt for each one *before* you build it. Build the smallest version that produces that receipt.

Week 2. Break each one on purpose. Did you find out? If not, fix the receipt – not the agent. Then, and only then, add the fourth.

That's it. That's the whole method. It is not sophisticated and it is the reason my systems are still standing.

THE RECEIPTS	github.com/kurenn/field-notes – every time one of my systems lied to me, written up in public, including the ones that made me look bad.
THE NUMBERS	kurenn.dev – two multi-agent architectures, ten real tasks, \$144.72 on the meter, and the finding that 69% of the bill was agents talking to each other.
THE TOOLS	github.com/kurenn – the agent teams, the audit tools and the prototyping kit I use for all of this. Open source, free.
SAY HELLO	x.com/kurenn · LinkedIn – tell me which three you picked. I read every one, and I'll tell you if I think you picked wrong.
